

# Kernel methods

Narayana Santhanam

EE 645  
Jan 27, 2026

## Linear → Non-linear

Two approaches:

Kernel methods

## Linear → Non-linear

Two approaches:

Kernel methods  
guarantees, well understood

# Linear → Non-linear

Two approaches:

Kernel methods

guarantees, well understood  
computationally intensive (small data)

Neural networks

less well understood

# Linear → Non-linear

Two approaches:

Kernel methods

guarantees, well understood  
computationally intensive (small data)

Neural networks

less well understood  
computationally cheap (large data)

## Next couple of weeks

Kernel methods

## Next couple of weeks

Kernel methods

High level picture

## Next couple of weeks

Kernel methods

High level picture

Support Vector classification and regression

## Next couple of weeks

Kernel methods

High level picture

Support Vector classification and regression

Kernel PCA

## Next couple of weeks

Kernel methods

High level picture

Support Vector classification and regression

Kernel PCA

Random Fourier Futures

## Next couple of weeks

Kernel methods

High level picture

Support Vector classification and regression

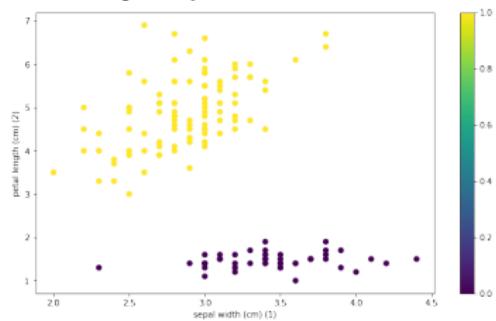
Kernel PCA

Random Fourier Features

Credit risk, Power data

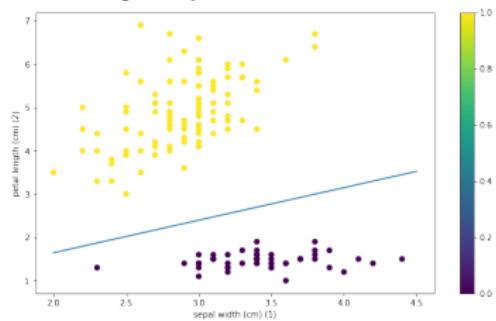
# Classification

Linearly separable



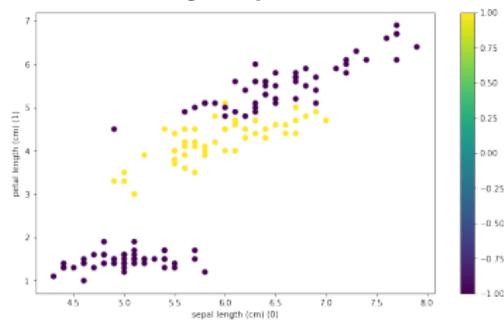
# Classification

Linearly separable



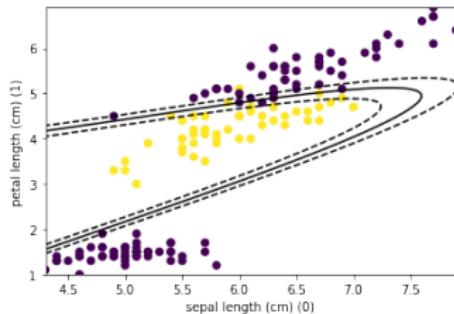
# Classification

Not linearly separable



# Classification

Not linearly separable



Boundaries not linear..

but are linear in a higher dimensional space!

## Closer look

Principled approach for linear → nonlinear

Powerful, yet generalizes well

Often explainable

at least more than other state of art

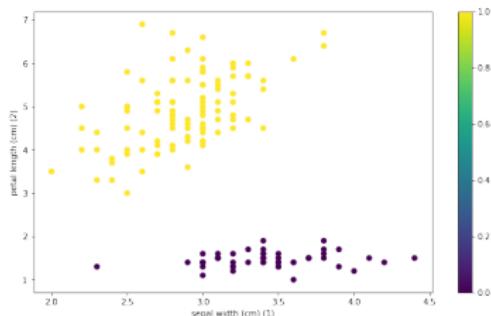
New advances increase reach (more data)

but not as much as NN

# Stepping back

Linearly separable

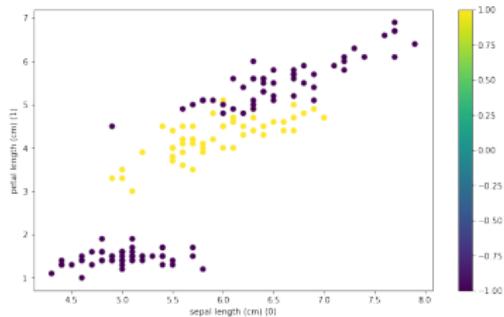
Where would you draw a *linear* classifier?



# Stepping back

Not linearly separable

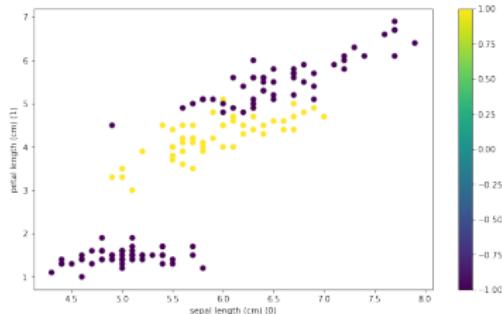
Where would you draw a *linear* classifier?



# Stepping back

Not linearly separable

Where would you draw a *linear* classifier?

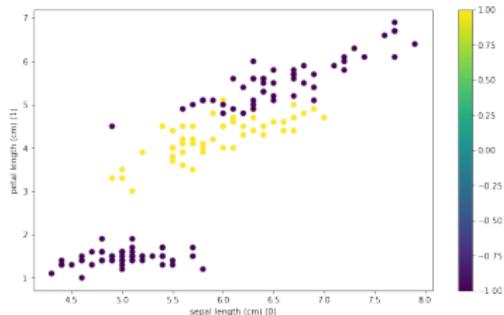


Minimum number of mistakes?

# Stepping back

Not linearly separable

Where would you draw a *linear* classifier?

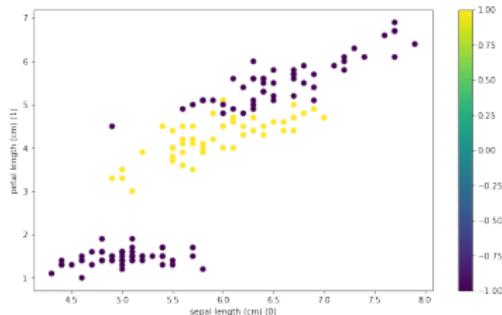


Minimum number of mistakes?  
infeasible, NP-hard

# Stepping back

Not linearly separable

Where would you draw a *linear* classifier?



Minimum number of mistakes?

infeasible, NP-hard

Instead: variation of  $\ell_1$  loss, sum of margins

## Setting up linear classification

Distance of  $\mathbf{z}$  from a plane  $\mathbf{w}^T \mathbf{x} - b = 0$ ?

## Setting up linear classification

Distance of  $\mathbf{z}$  from a plane  $\mathbf{w}^T \mathbf{x} - b = 0$ ?

$$\frac{\mathbf{w}^T \mathbf{z} - b}{\|\mathbf{w}\|}$$

# Formulation of Support Vector Classification

Analyze this in some detail

# Formulation of Support Vector Classification

Analyze this in some detail

One of the key advantages: interpretability

Comes through formulation

# Formulation of Support Vector Classification

Analyze this in some detail

One of the key advantages: interpretability

Comes through formulation

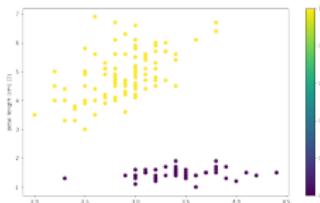
Primal/Dual formulation is a key optimization idea

accelerations of training neural networks

resource allocation/optimization in economics,

urban planning

# Formulation



Linearly separable points:  
Training data:

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)$$

Margin (closest distance from separating boundary)

$$\gamma(\mathbf{w}, b) = \min_{\mathbf{x}_i} \frac{|\mathbf{w}^T \mathbf{x}_i - b|}{\|\mathbf{w}\|}$$

Redundant parameterization scaling  $\mathbf{w}, b$  by any number does not change the margin

## Formulation

Support vector machine formulation:

$$\mathbf{w}^*, b^* = \arg \max_{\mathbf{w}, b} \gamma(\mathbf{w}, b)$$

subject to  $y_i(\mathbf{w}^T \mathbf{x}_i - b) \geq 0$  (for all training  $(\mathbf{x}_i, y_i)$ )

**Redundant parameterization:** scaling  $\mathbf{w}, b$  changes nothing

- only the directions/intercepts matter
- represent by scaling of  $\mathbf{w}, b$  that ensures

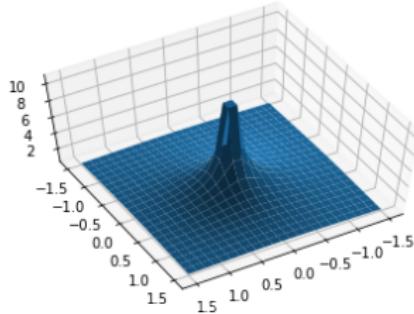
$$\min_{\mathbf{x}_i} |\mathbf{w}^T \mathbf{x}_i - b| = 1$$

# Formulation

Support vector machine formulation  
(removing the redundant formulation)

$$\mathbf{w}^*, b^* = \arg \max_{\mathbf{w}, b} \frac{1}{\|\mathbf{w}\|}$$

subject to  $y_i(\mathbf{w}^T \mathbf{x}_i - b) \geq 1$  (for all training  $(\mathbf{x}_i, y_i)$ )



## Formulation

$\frac{1}{\|\mathbf{w}\|}$  is convex, but...

*maximizing* (convex domain) isn't convex optimization  
minimize convex/maximize concave in convex domain

But it is easy to come with a convex formulation

$$\mathbf{w}^*, b^* = \arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

subject to  $y_i(\mathbf{w}^T \mathbf{x}_i - b) \geq 1$  (for all pairs  $(\mathbf{x}_i, y_i)$ )

# A bit on convex optimization

General concepts beyond support vector machine formulation

## A bit on convex optimization

General concepts beyond support vector machine formulation

Lagrangian:

$$L(\mathbf{w}, b, \Lambda) = \frac{1}{2} \|\mathbf{w}\|^2 + \sum_i \lambda_i (1 - y_i(\mathbf{w}^\top \mathbf{x}_i - b))$$

$L(\mathbf{w}, b, \Lambda)$  is a key idea in any constrained optimization

# Primal/dual formulation

Neat property of Lagrangians:

Optimal point

$$\min_{\mathbf{w}, b} \max_{\Lambda \geq 0} L(\mathbf{w}, b, \Lambda) \text{ (primal formulation)}$$

Dual formulation (Nash, von Neumann)

$$\max_{\Lambda \geq 0} \min_{\mathbf{w}, b} L(\mathbf{w}, b, \Lambda) \text{ (dual formulation)}$$

# Primal/dual formulation

Neat property of Lagrangians:

Optimal point

$$\min_{\mathbf{w}, b} \max_{\Lambda \geq 0} L(\mathbf{w}, b, \Lambda) \text{ (primal formulation)}$$

Dual formulation (Nash, von Neumann)

$$\max_{\Lambda \geq 0} \min_{\mathbf{w}, b} L(\mathbf{w}, b, \Lambda) \text{ (dual formulation)}$$

Incidentally Nash won both the Nobel Prize in Econ and the Abel Prize

# Primal/dual formulation

For free:

$$\min_{\mathbf{w}, b} \max_{\Lambda \geq 0} L(\mathbf{w}, b, \Lambda) \geq \max_{\Lambda \geq 0} \min_{\mathbf{w}, b} L(\mathbf{w}, b, \Lambda)$$

But equality in many cases

- in many convex formulations, including our current case
- so one could solve either version
- dual in kernel methods very insightful for explainability

## Dual formulation: 2 key insights

- Setting gradient of  $L(\mathbf{w}, b, \Lambda)$  to 0, optimal  $\mathbf{w}^*$  satisfies

$$\mathbf{w}^* = \sum_i \lambda_i y_i \mathbf{x}_i$$

**Representer theorem:** solution  $\mathbf{w}^*$  is linear combination of inputs

## Dual formulation: 2 key insights

- Setting gradient of  $L(\mathbf{w}, b, \Lambda)$  to 0, optimal  $\mathbf{w}^*$  satisfies

$$\mathbf{w}^* = \sum_i \lambda_i y_i \mathbf{x}_i$$

**Representer theorem:** solution  $\mathbf{w}^*$  is linear combination of inputs

- Finding  $\Lambda$ s

$$\max_{\Lambda \geq 0} \sum \lambda_i - \frac{1}{2} [\lambda_1 y_1 \quad \dots \quad \lambda_n y_n] X X^T \begin{bmatrix} \lambda_1 y_1 \\ \vdots \\ \lambda_n y_n \end{bmatrix}$$

Data only shows up through dot products ( $XX^T$ )  
Crux of Kernel approach to nonlinearity